

1 Spreadsheet Data Extractor (SDE): A Performance-Optimized, 2 User-Centric Tool for Transforming Semi-Structured Excel 3 Spreadsheets into Relational Data

4 ANONYMOUS AUTHOR(S)
5
6 SUBMISSION ID:

7 Spreadsheets are ubiquitous tools used across various domains. Despite their widespread use, analyzing and
8 utilizing data stored in spreadsheets poses significant challenges due to their semi-structured nature. Data
9 in spreadsheets are often formatted primarily for human readability, employing layouts and styles that are
10 easily understood by people but are difficult for automated systems to interpret. While these unstructured
11 formats offer advantages—such as providing an easily comprehensible hierarchy of metadata—they complicate
12 automated data extraction. This paper introduces the Spreadsheet Data Extractor (SDE), an open-source tool
13 designed to convert semi-structured spreadsheet data into structured formats without requiring programming
14 knowledge. Building upon previous work, we have enhanced the SDE with incremental loading of worksheets,
15 accurate rendering of cell dimensions by directly parsing the Excel file's XML content, and performance
16 optimizations to handle large datasets efficiently. We compare our tool with existing solutions and demonstrate
17 its effectiveness through performance evaluations, highlighting its potential to facilitate efficient and reliable
18 data extraction from diverse spreadsheet formats.

19 CCS Concepts: • Applied computing → Spreadsheets; • Information systems → Data cleaning;
20 • Software and its engineering → Extensible Markup Language (XML); • Human-centered computing
21 → Graphical user interfaces; • Theory of computation → Data compression.

22 Additional Key Words and Phrases: Spreadsheets, Data cleaning, Relational Data, Excel, XML, Graphical user
23 interfaces

24
25 **ACM Reference Format:**

26 Anonymous Author(s). 2025. Spreadsheet Data Extractor (SDE): A Performance-Optimized, User-Centric Tool
27 for Transforming Semi-Structured Excel Spreadsheets into Relational Data. In . ACM, New York, NY, USA,
28 16 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

29 **1 INTRODUCTION**

30 Spreadsheets are ubiquitous tools used across various domains, including healthcare [3], nonprofit
31 organizations [10, 13], finance, commerce, academia, and government [6]. Despite their widespread
32 use, analyzing and utilizing data stored in spreadsheets poses significant challenges for data
33 analysis and utilization, primarily due to their semi-structured nature. Data within spreadsheets
34 are frequently organized for human readability, featuring layouts with empty cells, merged cells,
35 hierarchical headers, and multiple tables. While these formats enhance comprehensibility for
36 users, they impede machine readability and automated data processing. This semi-structured
37 organization complicates the extraction of meaningful insights, especially when attempting to
38 integrate spreadsheet data into more robust and scalable data management systems.

39 The reliance on spreadsheets as ad-hoc solutions poses several limitations:

40 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee
41 provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the
42 full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored.
43 Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires
44 prior specific permission and/or a fee. Request permissions from permissions@acm.org.

45 *ACM PODS '25, June 22–27, 2025, Berlin, Germany*

46 © 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

47 ACM ISBN ...\$15.00

48 <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

- **Data Integrity and Consistency:** Spreadsheets are prone to errors, such as duplicate entries, inconsistent data formats, and inadvertent modifications, which can compromise data integrity. [? ?]
- **Scalability Issues:** As datasets grow in size and complexity, spreadsheets become less efficient for data storage and retrieval, leading to performance bottlenecks. [? ?]
- **Limited Query Capabilities:** Unlike databases, spreadsheets lack advanced querying and indexing features, restricting users from performing complex data analyses.

Transitioning from these ad-hoc spreadsheet solutions to standardized database systems offers numerous benefits:

- **Enhanced Data Integrity:** Databases enforce data validation rules and constraints, ensuring higher data quality and consistency.
- **Improved Scalability:** Databases are designed to handle large volumes of data efficiently, supporting complex queries and transactions without significant performance degradation.
- **Advanced Querying and Reporting:** Databases provide powerful querying languages like SQL, enabling sophisticated data analysis and reporting capabilities.
- **Seamless Integration:** Databases facilitate easier integration with various applications and services, promoting interoperability and data sharing across platforms.

Given the abundance of existing spreadsheet data and the clear advantages of database systems, there is a pressing need for tools that can bridge the gap between these two formats. Automated and accurate data extraction from spreadsheets into relational database formats is essential for organizations to leverage their data assets effectively.

Previous work by Aue et al. introduced a tool that facilitates data extraction from Excel files [1]. While effective, their solution faced performance issues and inaccuracies in rendering cell dimensions, limiting its usability with large and complex datasets.

In this paper, we aim to transform semi-structured spreadsheet data into machine-readable formats by building upon their work. We present the Spreadsheet Data Extractor (SDE), an enhanced tool that enables users to define data hierarchies through cell selection without any programming knowledge. [?] Our enhancements address key limitations of the existing solution, making data extraction from complex spreadsheets more efficient and user-friendly.

1.1 Contributions

Our main contributions are as follows:

- (1) We release the SDE under the open-source GNU General Public License v3.0, promoting community access and collaboration. [?]
- (2) We implement incremental loading of worksheets to enhance performance, allowing the tool to handle large Excel files efficiently.
- (3) We accurately render row heights and column widths by parsing XML data, ensuring that the spreadsheet's visual representation closely matches that of Excel.
- (4) We optimize the rendering engine to draw only the visible cells, significantly improving performance when dealing with large datasets.
- (5) We integrate the selection hierarchy, worksheet view, and output preview into a unified interface, streamlining the user experience.

2 RELATED WORK

The extraction of relational data from semi-structured documents, particularly spreadsheets, has garnered significant attention due to their ubiquitous use across domains such as business, government, and scientific research. Several frameworks and tools have been developed to address

99 the challenges of converting flexible spreadsheet formats into normalized relational forms suitable
100 for data analysis and integration. Notable among these are **DeExcelerator** [7], **XLIndy** [8],
101 **FLASHRELATE** [2], **Senbazuru** [4], **TableSense** [5], and the approach by Aue et al. [1], on which
102 our work builds.

103 2.1 Aue et al.'s Converter

104 Aue et al. [1] developed a tool to facilitate data extraction from Excel spreadsheets by leveraging
105 the Dart excel package [?] to process .xlsx files. This tool allows users to define data hierarchies
106 by selecting relevant cells containing data and metadata. However, the approach faced significant
107 performance bottlenecks due to the excel package's requirement to load the entire .xlsx file into
108 memory, resulting in slow response times, particularly for large files

109 In addition to memory issues, the tool calculated row heights and column widths based solely
110 on cell content, ignoring the dimensions specified in the original Excel file. This led to rendering
111 discrepancies between the tool and the original spreadsheet. Furthermore, the tool rendered all
112 cells, regardless of their visibility within the viewport, significantly degrading performance when
113 handling worksheets with large numbers of cells.

114 2.2 DeExcelerator

115 Eberius et al. [7] introduced **DeExcelerator**, a framework that transforms partially structured
116 spreadsheets into first normal form relational tables using heuristic-based extraction phases. It
117 addresses challenges such as table detection, metadata extraction, and layout normalization. While
118 effective in automating normalization, its reliance on predefined heuristics limits adaptability to
119 heterogeneous or unconventional spreadsheet formats, highlighting the need for more flexible
120 approaches.

121 2.3 XLIndy

122 Koci et al. [8] developed **XLIndy**, an interactive Excel add-in with a Python-based machine learning
123 backend. Unlike DeExcelerator's fully automated heuristic approach, XLIndy integrates machine
124 learning techniques for layout inference and table recognition, enabling a more adaptable and
125 accurate extraction process. XLIndy's interactive interface allows users to visually inspect extraction
126 results, adjust configurations, and compare different extraction runs, facilitating iterative fine-tuning.
127 Additionally, users can manually revise predicted layouts and tables, saving these revisions as
128 annotations to improve classifier performance through (re)-training. This user-centric approach
129 enhances the tool's flexibility, allowing it to accommodate diverse spreadsheet formats and user-
130 specific requirements more effectively than purely heuristic-based systems.

131 2.4 FLASHRELATE

132 Barowy et al. [2] presented **FLASHRELATE**, an approach that empowers users to extract structured
133 relational data from semi-structured spreadsheets without requiring programming expertise.
134 FLASHRELATE introduces a domain-specific language, **FLARE**, which extends traditional regular
135 expressions with spatial constraints to capture the geometric relationships inherent in spreadsheet
136 layouts. Additionally, FLASHRELATE employs an algorithm that synthesizes FLARE programs
137 from a small number of user-provided positive and negative examples, significantly simplifying the
138 automated data extraction process.

139 FLASHRELATE distinguishes itself from both DeExcelerator and XLIndy by leveraging programming-
140 by-example (PBE) techniques. While DeExcelerator relies on predefined heuristic rules and XLIndy
141 incorporates machine learning models requiring user interaction for fine-tuning, FLASHRELATE
142 allows non-expert users to define extraction patterns through intuitive examples. This approach

148 lowers the barrier to entry for extracting relational data from complex spreadsheet encodings,
149 making the tool accessible to a broader range of users.

150 151 2.5 Senbazuru

152 Chen et al. [4] introduced **Senbazuru**, a prototype Spreadsheet Database Management System
153 (SSDBMS) designed to extract relational information from a large corpus of spreadsheets. Senbazuru
154 addresses the critical issue of integrating data across multiple spreadsheets, which often lack explicit
155 relational metadata, thereby hindering the use of traditional relational tools for data integration
156 and analysis.

157 Senbazuru comprises three primary functional components:

158 (1) **Search**: Utilizing a textual search-and-rank interface, Senbazuru enables users to quickly
159 locate relevant spreadsheets within a vast corpus. The search component indexes spread-
160 sheets using Apache Lucene, allowing for efficient retrieval based on relevance to user
161 queries.

162 (2) **Extract**: The extraction pipeline in Senbazuru consists of several stages:

- 163 • **Frame Finder**: Identifies data frame structures within spreadsheets using Conditional
164 Random Fields (CRFs) to assign semantic labels to non-empty rows, effectively detecting
165 rectangular value regions and associated attribute regions.

- 166 • **Hierarchy Extractor**: Recovers attribute hierarchies for both left and top attribute
167 regions. This stage also incorporates a user-interactive repair interface, allowing users
168 to manually correct extraction errors, which the system then generalizes to similar
169 instances using probabilistic methods.

- 170 • **Tuple Builder and Relation Constructor**: Generates relational tuples from the
171 extracted data frames and assembles these tuples into coherent relational tables by
172 clustering attributes and recovering column labels using external schema repositories
173 like Freebase and YAGO.

174 (3) **Query**: Supports basic relational operations such as selection and join on the extracted
175 relational tables, enabling users to perform complex data analysis tasks without needing to
176 write SQL queries.

177 Senbazuru's ability to handle hierarchical spreadsheets, where attributes may span multiple
178 rows or columns without explicit labeling, sets it apart from earlier systems like DeExcelerator and
179 XLIIndy. By employing machine learning techniques and providing user-friendly repair interfaces,
180 Senbazuru ensures high-quality extraction and facilitates the integration of spreadsheet data into
181 relational databases.

182 183 2.6 TableSense

184 Dong et al. [5] developed **TableSense**, an end-to-end framework for spreadsheet table detection
185 using Convolutional Neural Networks (CNNs). TableSense addresses the diversity of table structures
186 and layouts by introducing a comprehensive cell featurization scheme, a Precise Bounding Box
187 Regression (PBR) module for accurate boundary detection, and an active learning framework to
188 efficiently build a robust training dataset.

189 While **DeExcelerator**, **XLIIndy**, **FLASHRELATE**, and **Senbazuru** focus primarily on transforming
190 spreadsheet data into relational forms through heuristic, machine learning, and programming-
191 by-example approaches, **TableSense** specifically targets the accurate detection of table boundaries
192 within spreadsheets using deep learning techniques. Unlike region-growth-based methods em-
193 ployed in commodity spreadsheet tools, which often fail on complex table layouts, TableSense
194 achieves superior precision and recall by leveraging CNNs tailored for the unique characteristics of
195

197 spreadsheet data. However, TableSense focuses on table detection and visualization, allowing users
198 to generate diagrams from the detected tables but does not provide functionality for exporting the
199 extracted data for further analysis.

200 201 2.7 Comparison and Positioning

202 While **DeExcelerator**, **XLIndy**, **FLASHRELATE**, **Senbazuru**, and **TableSense** each offer unique
203 approaches to spreadsheet data extraction, they share certain limitations. Many of these tools
204 are not readily accessible: **FLASHRELATE** and **TableSense** are proprietary, and **Senbazuru**,
205 **XLIndy**, and **DeExcelerator** are discontinued projects with limited or no source code availability.
206 In contrast, we contribute our spreadsheet data extractor under the GNU General Public License
207 v3.0, allowing the community to access, use, and improve the tool freely.

208 Moreover, unlike the aforementioned tools that rely on heuristics, machine learning, or AI
209 techniques—which can introduce errors requiring users to identify and correct—we adopt a user-
210 centric approach that gives users full control over data selection and metadata hierarchy definition.
211 While this requires more manual input, it eliminates the uncertainty and potential inaccuracies
212 associated with automated methods. To streamline the process and enhance efficiency, our tool
213 includes user-friendly features such as the ability to duplicate hierarchies of columns and tables,
214 and to move them over similar structures for reuse, reducing the need for repetitive configurations.

215 By combining the strengths of manual control with enhanced user interface features and per-
216 formance optimizations, our tool offers a robust and accessible solution for extracting relational
217 data from complex and visually intricate spreadsheets. These enhancements not only improve
218 performance and accuracy but also elevate the overall user experience, making our tool a valuable
219 asset for efficient and reliable data extraction from diverse spreadsheet formats.

220 221 3 METHODOLOGY

222 In this section, we detail the design and implementation of the Spreadsheet Data Extractor (SDE),
223 emphasizing its user-centric approach and performance optimizations. The SDE enables users to
224 transform semi-structured spreadsheet data into structured, machine-readable formats without
225 requiring programming expertise. We achieve this through an intuitive interface that allows for
226 cell selection and hierarchy definition, incremental loading of worksheets, accurate rendering of
227 cell dimensions, and optimized performance for handling large datasets by incrementally loading
228 worksheets and by rendering only the cells that are currently visible in the view.

229 230 3.1 User-Centric Data Extraction

231 The core functionality of the SDE revolves around allowing users to select cells containing data
232 and metadata to define a data hierarchy. This process is facilitated through a graphical interface
233 that displays the spreadsheet and allows for intuitive selection and manipulation of the selection
234 hierarchy.

235 3.1.1 *Hierarchy Definition.* Users can select individual cells or ranges of cells by clicking and using
236 shift-click for multi-selection. These selections represent either data or metadata.

237 The selected cells are organized into a hierarchical tree structure, where each node represents a
238 data element, and child nodes represent nested data or metadata. This hierarchy defines how the
239 data will be transformed into a structured format.

240 3.1.2 *Reusability and Efficiency.* To optimize the extraction process and reduce repetitive tasks,
241 the SDE allows users to duplicate previously defined hierarchies and apply them to similar regions
242 within the spreadsheet. This feature is particularly useful for spreadsheets with repeating structures,
243 such as multiple tables with the same format.

246 3.2 Example Workflow

247 Consider a spreadsheet containing statistical forecasts of future nursing staff availability in Germany
 248 [11]. Figure 1 shows the SDE interface, which consists of three main components:

249 **Hierarchy Panel (Top Left):** Displays the hierarchy of cell selections, initially empty.

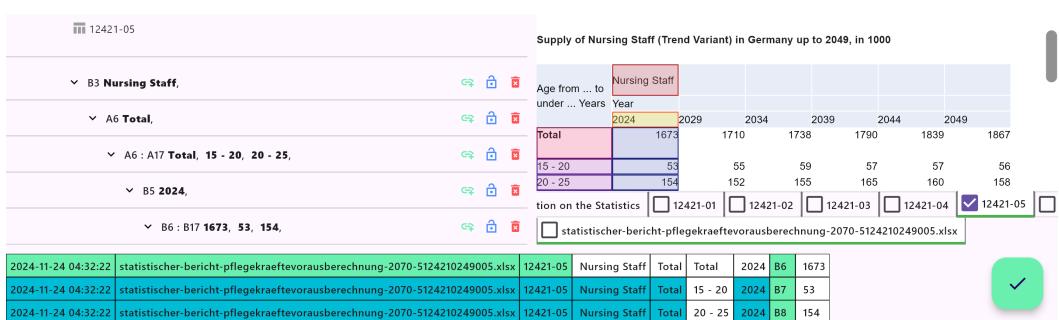
250 **Spreadsheet View (Top Right):** Shows the currently opened Excel file and the currently selected
 251 worksheet for cell selection.

252 **Output Preview (Bottom):** Provides immediate feedback on the data extraction based on
 253 current selections.



264 Fig. 1. The SDE Interface Overview.
 265
 266

267 **3.2.1 Selection of the First Column.** The user adds a node to the hierarchy and selects the cell
 268 containing the metadata "Nursing Staff" (Figure 2). This cell represents metadata that is common to
 269 all cells in this worksheet. Therefore, it should be selected first and should appear at the beginning
 270 of each row in the output CSV file.



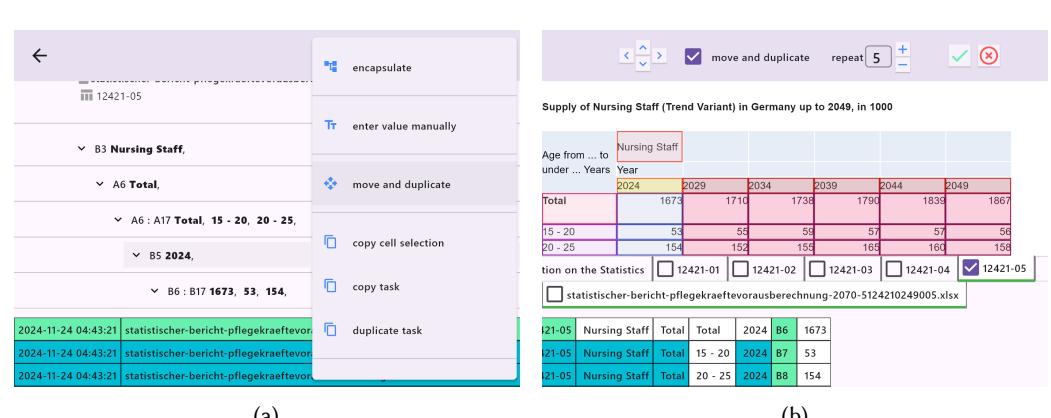
282 Fig. 2. Selection of the First Column Metadata
 283
 284

285 Within this node, the user adds a child node and selects the cell "Total" which serves both as a
 286 table header and a row label. This selection represents the table header of the first subtable. The
 287 user adds another child node and selects the range of cells containing row labels ("Total," "15-20",
 288 "20-25", ...) by clicking the first cell and shift-clicking the last cell. A child node is added under the
 289 row labels node, and the user selects the year "2024" and after that another child node is added
 290 under the year node and the user selects the corresponding data cells (e.g., "1673," "53", "154", ...).
 291 The hierarchy is now filled with 5 nodes, each having an embedded child node except for the last one.
 292 In the upper right area, the selected cells are displayed with different colors for each node. In
 293 the lower area, the output is displayed. For each node that is embedded as a child, another column
 294

295 appears in the output. If multiple cells are selected, the values from the list also appear as values in
 296 a new row in the output.

297 3.2.2 *Duplication of the Column Hierarchy.* To avoid repetitive manual entry for additional years,
 298 the user duplicates the hierarchy for "2024" and adjusts the cell selections to include data for
 299 subsequent years (e.g., "2025," "2026") using the "Move and Duplicate" feature.
 300

301 To do this, the user selects the node of the first column "2024" and right-clicks on it. A popup
 302 opens in which the action "move and duplicate" appears, which should then be clicked, as shown
 303 in Figure 3a.



317 Fig. 3. Move and Duplicate Feature in Action

318 Subsequently, a series of buttons opens in the app bar at the top right, allowing the user to move
 319 the cell selections of the node as well as all child nodes, as seen in Figure 3b. By pressing the button
 320 to move the selection by one unit to the right, the next column is selected; however, this would
 321 also deselect the first column since the selection was moved. To preserve the first column, the
 322 "move and duplicate" checkbox can be activated. This creates the shifted selection in addition to the
 323 original selection. However, the changes are only applied when the accept button is clicked. The
 324 next columns could also be selected in the same way. But this can be done faster, because instead
 325 of moving the selection and duplicating it only once, the "repeat" input field can be filled with as
 326 many repetitions as there are columns. By entering the number 5, the selection of the first column
 327 is shifted 5 times by one unit to the right and duplicated at each step.

328 The user reviews the selections in the spreadsheet view, where each selection is highlighted in a
 329 different color corresponding to its node in the hierarchy. Only after the user has reviewed the
 330 shifted and duplicated selections in the worksheet and clicked the accept button are the nodes in
 331 the hierarchy created as desired. We present the results of these runtime measurements in Figure 4.

332 The user reviews the selections in the spreadsheet view, where each selection is highlighted in a
 333 different color corresponding to its node in the hierarchy. Erst nachdem der Nutzer die verschobenen
 334 und duplizierten selektionen in der Worksheet ansieht überprüft hat und den Akzeptieren button
 335 geklickt hat, werden die Knoten in der Hierachie wie gewünscht angelegt. Das Ergebnis dieser
 336 Operation ist in Abbildung 4 zu sehen.

337 3.2.3 *Duplicating the Table Hierarchy.* The same method that worked effectively for duplicating
 338 the columns can now be applied to the subtables, as shown in Figure 5.

344 12421-05

345 Supply of Nursing Staff (Trend Variant) in Germany up to 2049, in 1000

346 ✓ B3 Nursing Staff,

347 ✓ A6 Total,

348 ✓ A6 : A17 Total, 15 - 20, 20 - 25,

349 ✓ B5 2024,

350 ✓ B6 : B17 1673, 53, 154,

351 2024-11-24 04:47:51 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total Total 2024 B6 1673

352 2024-11-24 04:47:51 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 15 - 20 2024 B7 53

353 2024-11-24 04:47:51 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 20 - 25 2024 B8 154

Fig. 4. Resulting Hierarchy After Move and Accept

354

355

356

357 ←

358 ✓ B3 Nursing Staff,

359 ✓ A6 Total,

360 ✓ A6 : A17 Total, 15 - 20, 20 - 25,

361 ✓ B5 2024,

362 ✓ B6 : B17 1673, 53, 154,

363 ✓ C5 2029,

364 ✓ C6 : C17 1710, 55, 152,

365 ✓ D5 2034,

366 ✓ D6 : D17 1738, 59, 155,

367 ✓ E5 2039,

368 ✓ E6 : E17 1790, 57, 165,

369 ✓ F5 2044,

370 ✓ F6 : F17 1839, 57, 160,

371 ✓ G5 2049,

372 ✓ G6 : G17 1867, 56, 158,

373 2024-11-24 04:52:37 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total Total 2024 B6 1673

374 2024-11-24 04:52:37 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 15 - 20 2024 B7 53

375 2024-11-24 04:52:37 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 20 - 25 2024 B8 154

376 2024-11-24 04:52:37 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 25 - 30 2024 B9 173

377 2024-11-24 04:52:37 statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.xlsx 12421-05 Nursing Staff Total 30 - 35 2024 B10 178

Fig. 5. Selection of All Cells in the Subtables by Duplicating the Hierarchy of the First Table

378 By selecting the node with the value "Total" and clicking the "Move and Duplicate" button, we
 379 can apply the selection of the "Total" subtable to the other subtables. This involves shifting the
 380 table downward by as many rows as necessary to overlap with the subtable below.

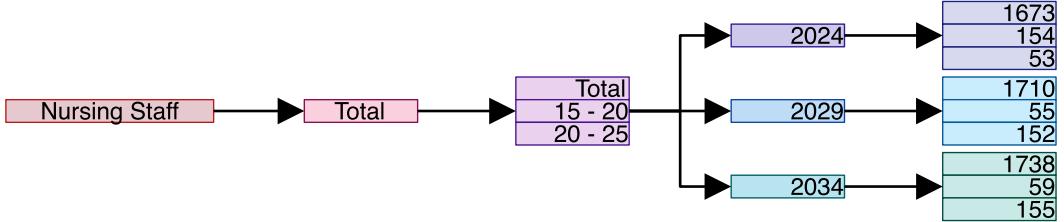
381 However, there is a minor issue: the child nodes of the "Total" node also include the column
 382 headers. If these column headers were repeated in the subtables below, shifting the selections
 383 downward would work without modification. Since these cells are not repeated in the subtables,
 384 we need to prevent the column headers cells from moving during the duplication process.

385 To achieve this, we can exclude individual nodes from being moved by locking their selection.
 386 This is done by clicking the padlock icon on the corresponding nodes, which freezes their cell
 387 selection and keeps them fixed at their original position, regardless of other cells being moved.

388 Therefore, we identify and select the nodes containing the column headers—specifically, the
 389 years 2024 to 2049—and lock their selection using the padlock button. By shifting the selection
 390 downward and duplicating it, we can easily move and duplicate the cell selections for the subtables
 391 below. By setting the number of repetitions to 2, all subtables are completely selected.

393 3.3 Cross Product Transformation

394 The graph resulting from the selected hierarchy is shown in Figure 6. To simplify the explanation,
 395 the example is limited to the first three columns and the first three rows of the first sub-table.
 396

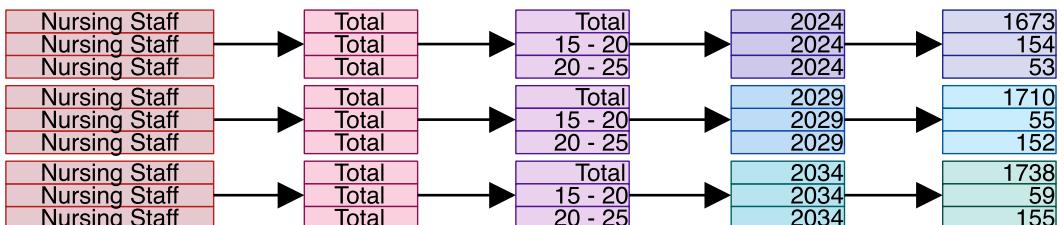


405 Fig. 6. Illustration of the Cross Product Transformation Before Application

406 Once the hierarchy is defined, the **Structured Data Engine (SDE)** applies a cross-product
 407 transformation to generate a relational format from the selection graph. This transformation
 408 consists of two key steps:

- 410 (1) **Node Duplication:** Nodes with multiple incoming or outgoing edges (e.g., the row labels
 411 node with the values *Total*, *15-20*, *20-25*) are duplicated to ensure that each edge connects to a
 412 unique instance of the node. This adjustment replaces the original many-to-one relationships
 413 with one-to-one mappings for each edge.
- 414 (2) **Value Replication:** Nodes containing single values (e.g., the year *2024*) are replicated
 415 to align with the number of values associated with the connected node. This ensures a
 416 consistent structure in the relational output, maintaining alignment across all hierarchical
 417 levels.

418 The resulting graph after the cross-product transformation is shown in Figure 7. With this
 419 transformation applied to the selected hierarchy, the SDE generates a structured output that reflects
 420 the original spreadsheet's hierarchical relationships, enabling users to analyze and integrate the
 421 data effectively.



431 Fig. 7. Illustration of the Cross Product Transformation After Application

434 3.4 Incremental Loading of Worksheets

435 Opening large Excel files traditionally involves loading the entire file and all its worksheets into
 436 memory before displaying any content. In files containing very large worksheets, this process
 437 can take several seconds to minutes, causing significant delays for users who need to access data
 438 quickly.

439 To facilitate efficient data extraction from multiple Excel files, we implemented a mechanism
 440 for incremental loading of worksheets within the SDE. Excel files (.xlsx format) are ZIP archives

442 containing a collection of XML files that describe the worksheets, styles, and shared strings. Key
 443 components include:

- 444 • **xl/sharedStrings.xml**: Contains all the unique strings used across worksheets, reducing
 445 redundancy.
- 446 • **xl/styles.xml**: Defines the formatting styles for cells, including fonts, colors, and borders.
- 447 • **xl/worksheets/sheetX.xml**: Represents individual worksheets (`sheet1.xml`, `sheet2.xml`,
 448 etc.).

449 Our solution opens the Excel file as a ZIP archive and initially extracts only the essential metadata
 450 and shared resources required for the application to function. This initial extraction includes:
 451

452 (1) **Metadata Extraction:**

453 We read the archive's directory to identify the contained files without decompressing them
 454 fully. This step is quick, taking only a few milliseconds, and provides information about the
 455 available worksheets and shared resources.

456 (2) **Selective Extraction:**

457 We immediately extract `sharedStrings.xml` and `styles.xml` because these files are small
 458 and contain information necessary for rendering cell content and styles across all worksheets.
 459 These files are parsed and stored in memory for quick access during rendering.

460 (3) **Deferred Worksheet Loading:**

461 The individual worksheet files (`sheetX.xml`) remain compressed and are loaded into mem-
 462 ory in their binary unextracted form. They are not decompressed or parsed at this stage.

463 (4) **On-Demand Parsing:**

464 When a user accesses a specific worksheet—either by selecting it in the interface or when a
 465 unit test requires data from it—the corresponding `sheetX.xml` file is then decompressed
 466 and parsed. This parsing occurs in the background and is triggered only by direct user
 467 action or programmatic access to the worksheet's data.

468 (5) **Memory Release:**

469 After a worksheet has been decompressed and its XML parsed, we release the memory
 470 resources associated with the parsed data. This approach prevents excessive memory usage
 471 and ensures that the application remains responsive even when working with multiple
 472 large worksheets.

473 By adopting this incremental loading approach, users experience minimal wait times when
 474 opening an Excel file. The initial loading is nearly instantaneous, allowing users to begin interacting
 475 with the application without delay. This contrasts with traditional methods that require loading all
 476 worksheets upfront, leading to significant wait times for large files.
 477

478 3.5 Rendering of Worksheets

479 To ensure that users can navigate worksheets without difficulty, we prioritize displaying the
 480 worksheets in a manner that closely resembles their appearance in Excel. This involves accurately
 481 rendering cell dimensions, formatting, and text behaviors.
 482

483 *3.5.1 Displaying Row Heights and Column Widths.* Our solution extracts information about column
 484 widths and row heights directly from the Excel file's XML structure. Specifically, we retrieve the
 485 column widths from the `width` attribute of the `<col>` elements and the row heights from the `ht`
 486 attribute of the `<row>` elements in the `sheetX.xml` files.

487 In Excel, column widths and row heights use units that do not directly map to pixels, requiring
 488 conversion for accurate on-screen rendering. Different scaling factors are needed for columns and
 489 rows. Through empirical testing, we derived the following scaling factors:
 490

- 491 • **Column Widths:** Multiply the *width* attribute by 7.
 492 • **Row Heights:** Multiply the *ht* attribute by $\frac{4}{3}$.

493 Despite research, we could not find official documentation explaining the rationale behind these
 494 specific scaling factors. This lack of documentation poses a challenge for accurately replicating
 495 Excel's rendering.

496 3.5.2 *Cell Formatting*. Cell formatting plays a crucial role in accurately representing the appearance
 497 of worksheets. Formatting information is stored in the *styles.xml* file, where styles are defined
 498 and later referenced in the *sheetX.xml* files as shown in Figure 8.

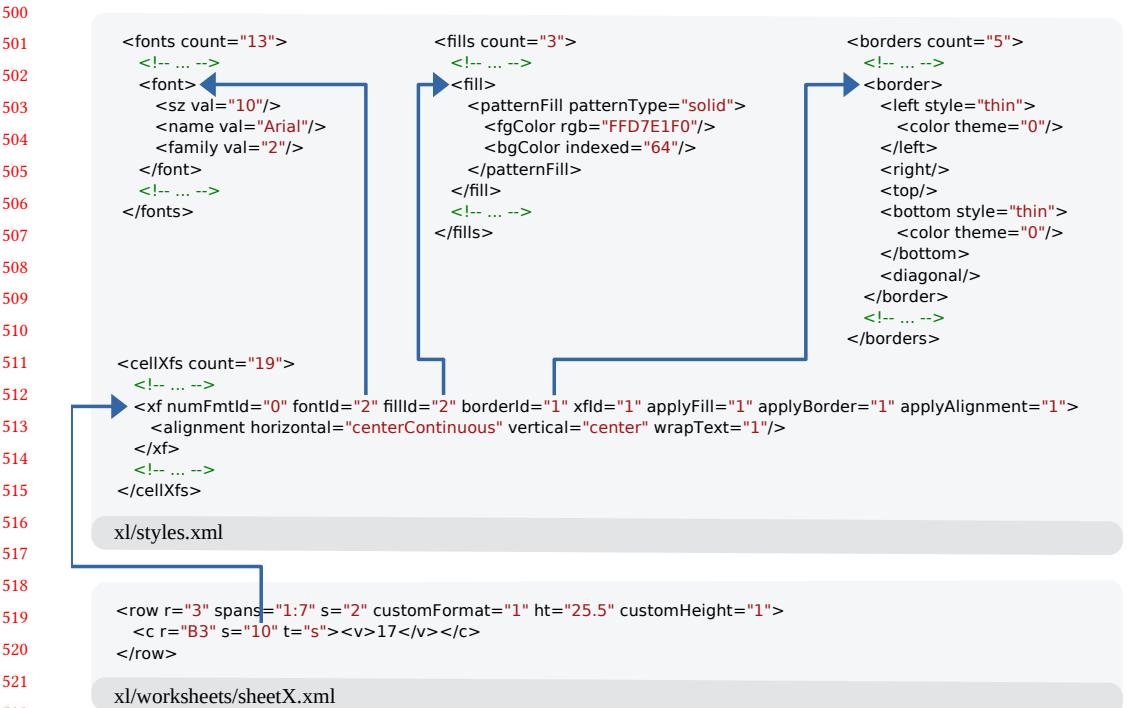


Fig. 8

525 Each cell in the worksheet references a style index through the *s* attribute, which points to the
 526 corresponding *xf* element within the *cellXfs* collection. These *xf* elements contain attributes
 527 such as *fontId*, *fillId*, and *borderId*, which reference specific font, fill (background), and border
 528 definitions located in the *fonts*, *fills*, and *borders* collections, respectively. By parsing these
 529 references, we can accurately apply the appropriate fonts, background colors, and border styles to
 530 each cell.

531 Through meticulous parsing and application of these formatting details, we ensure that the
 532 rendered worksheet closely mirrors the original Excel file, preserving the visual cues and aesthetics
 533 that users expect.

535 3.5.3 *Handling Text Overflow*. In Excel, when the content of a cell exceeds its width, the text may
 536 overflow into adjacent empty cells, provided those cells do not contain any data. If adjacent cells
 537 are occupied, Excel truncates the overflowing text at the cell boundary. Replicating this behavior is
 538 essential for accurate rendering and user familiarity.

We implemented text overflow handling by checking if the adjacent cell to the right is empty before allowing text to overflow. If the adjacent cell is empty, we extend the text rendering. If the adjacent cell contains data, we truncate the text at the boundary of the original cell.

Figure 1 illustrates this behavior. The text "Supply of Nursing Staff ..." extends into the neighboring cell because it is empty. If not for this handling, the text would be truncated at the cell boundary, leading to incomplete data display as shown in Figure 9.

[To the table of contents](#)

Supply of Nurs

Age from ... to under ... Years	Nursing Staff							
	Year	2024	2029	2034	2039	2044	2049	1839
		Total	1673	1710	1738	1790	2044	2049
2024	1673	1710	1738	1790	2044	2049	1839	1867

Fig. 9. Falsche Darstellung ohne overflow von zellen mit angrenzenden zellen ohne inhalt

By accurately handling text overflow, we improve readability and maintain consistency with Excel's user interface, which is crucial for users transitioning between Excel and our tool.

3.6 Performance Optimization

To ensure high frame rates even with large worksheets, we optimized the Spreadsheet Data Extractor to render only the cells that are currently visible to the user. Rendering the entire worksheet, especially when it contains thousands of cells, can significantly degrade performance. By focusing on the visible cells within the viewport, we reduce computational overhead and improve responsiveness.

We achieve this optimization by utilizing the `two_dimensional_scrollables` package [9]. This package provides functionality for efficiently handling two-dimensional scrolling regions, making it suitable for rendering large grids like spreadsheets.

Since we extract all column widths and row heights from the XML files, we can calculate the exact dimensions and positions of each cell. By accumulating the widths and heights, we determine the coordinates of each cell within the worksheet grid. These coordinates are essential for identifying which cells fall within the current viewport and should be rendered.

We consider the current scroll offsets along the horizontal (x -axis) and vertical (y -axis) directions. The viewport is defined by the width and height of the panel displaying the Excel worksheet. To determine the visible cells, we perform the following steps:

- **Horizontal Visibility:**

- Sum the column widths until the accumulated sum reaches the left edge of the viewport.
All cells to the left are ignored.
- Continue adding column widths until the sum exceeds the right edge of the viewport.
Cells beyond this point are also ignored.

- **Vertical Visibility:**

- 589 – Sum the row heights until the accumulated sum reaches the top edge of the viewport.
 590 All cells above are ignored.
 591 – Continue adding row heights until the sum surpasses the bottom edge of the viewport.
 592 Cells below this point are ignored.

593 By rendering only the cells within these boundaries, we significantly reduce the number of cells
 594 processed at any given time.

595 The two_dimensional_scrollables package provides interfaces where the logic for laying out
 596 the cells can be implemented. Parameters that describe the viewport, including the horizontal and
 597 vertical offsets as well as the viewport height and width, are supplied by these interfaces.

598 Algorithm 1 outlines our implementation of the overridden layoutChildSequence function,
 599 which is invoked by the two_dimensional_scrollables package to calculate and arrange the
 600 visible cells within the viewport.

603 Algorithm 1 Layout of Visible Spreadsheet Cells

```

604 1: Initialize Indices
605 2:   leadingColumnIndex ← column index corresponding to horizontalOffset
606 3:   leadingRowIndex ← row index corresponding to verticalOffset
607 4:   trailingColumnIndex ← column index corresponding to horizontalOffset + viewportWidth
608 5:   trailingRowIndex ← row index corresponding to verticalOffset + viewportHeight
609 6: Calculate Initial Offsets
610 7:   leadingColumnOffset ← sum of widths from the first column up to leadingColumnIndex
611 8:   leadingRowOffset ← sum of heights from the first row up to leadingRowIndex
612 9:   horizontalLayoutOffset ← leadingColumnOffset - horizontalOffset
613 10: for each columnIndex from leadingColumnIndex to trailingColumnIndex do
614 11:   verticalLayoutOffset ← leadingRowOffset - verticalOffset
615 12:   for each rowIndex from leadingRowIndex to trailingRowIndex do
616 13:     cell ← build or retrieve the cell at (columnIndex, rowIndex)
617 14:     Layout cell at position (horizontalLayoutOffset, verticalLayoutOffset)
618 15:     if a custom height is defined for row rowIndex then
619 16:       verticalLayoutOffset ← verticalLayoutOffset + height of row rowIndex
620 17:     else
621 18:       verticalLayoutOffset ← verticalLayoutOffset + defaultRowHeight
622 19:     end if
623 20:   end for
624 21:   columnWidth ← width of column columnIndex
625 22:   horizontalLayoutOffset ← horizontalLayoutOffset + columnWidth
626 23: end for
627
628

```

629 By applying this method, we render only the cells necessary for the current view, thereby
 630 optimizing performance and ensuring smooth user interactions even with large and complex
 631 worksheets.

632 4 EVALUATION

634 The fundamental approach of the Spreadsheet Data Extractor, based on the converter by Aue et
 635 al. [1], upon which we build, remains unchanged. The effectiveness of this approach has already
 636 been investigated. Aue et al. evaluated the extraction of data from over 500 Excel files. The time

required for each file was determined from a sample of 331 processed Excel files comprising 3,093 worksheets. On average, student assistants needed 15 minutes per file and 95 seconds per worksheet.

Our focus is on improving the user experience and optimizing the performance of the Spreadsheet Data Extractor. We enhanced the user experience by displaying the Excel worksheets similarly to how they appear in Excel and by reducing the number of required user interactions through the integration of the selection hierarchy, worksheet view, and output preview into a single interface. Performance was further improved by implementing incremental loading of Excel files and rendering only the visible cells.

4.1 Acceleration When Opening Files

To evaluate the performance improvements when opening Excel files, we conducted a series of tests using a large Excel file. We downloaded the entire collection of Excel files from the German Federal Statistical Office (Destatis)¹, which provides extensive statistical data across various domains. From this collection, we identified the largest file [12], which has a compressed size of 87 MB and an uncompressed size of 911 MB.

We performed three sets of tests to compare the performance of different methods for opening and reading data from this Excel file:

- (1) **VBA Script in Excel:** We wrote a VBA script that opens the Excel file and reads specific values from a worksheet. This script simulates how users might interact with the file using Excel's built-in capabilities. We ran this script 10 times, measuring the time required to open the file and read the values in each iteration.
- (2) **Spreadsheet Data Extractor:** We developed an equivalent unit test using the functions implemented in our Spreadsheet Data Extractor to open the same file and read the same cells. This test aimed to assess the performance of our tool under the same conditions. We repeated this test 10 times, recording the runtime for each execution.
- (3) **excel Package:** We tested the `excel` package [?] used by Aue et al. [1], creating a unit test that attempts to read the same cells from the file. This test was intended to benchmark the performance of the prior solution. However, due to limitations with the package, we encountered an out-of-memory exception during the second run; therefore, only the time from the first run was recorded.

All tests were conducted on a machine with the following specifications: Intel Core i5-10210U quad-core CPU at 1.60 GHz, 16 GB RAM, and a solid-state drive (SSD), running Windows 10 Pro. The version of Microsoft Excel used was Microsoft Office LTSC Professional Plus 2021.

The results of these runtime measurements are presented in Figure 10.

The Spreadsheet Data Extractor opened the worksheet with a median time of 120 milliseconds and an average time of 178 milliseconds. The first run took 668 milliseconds, possibly because the file was not yet cached and had to be loaded from the disk.

In contrast, Excel opened the worksheet with a median time of 40.281 seconds and an average time of 41.138 seconds.

The Dart `excel` package [?] used in the previous work took 13 minutes and 15 seconds to open the worksheet in the first run. The other nine runs could not be completed because an out-of-memory exception was thrown during the second run.

These results demonstrate that the Spreadsheet Data Extractor opens worksheets over two orders of magnitude faster than Excel and nearly four orders of magnitude faster than the Dart `excel` package used in prior work.

¹<https://www.destatis.de>

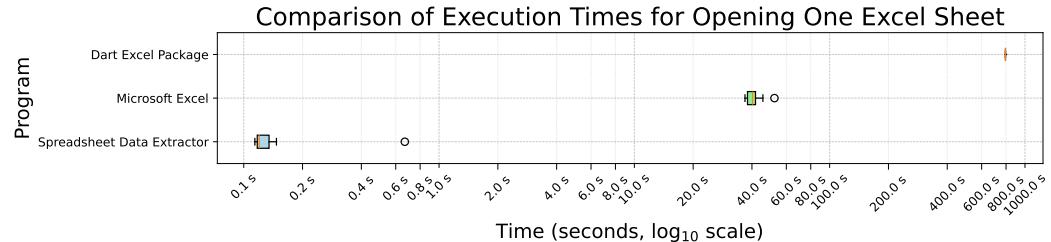


Fig. 10. Boxplot of Worksheet Opening Times Comparing the Spreadsheet Data Extractor, Microsoft Excel, and the Dart excel Package

5 FUTURE WORK

We plan to continue improving the Spreadsheet Data Extractor by implementing new features that enhance the user experience and address current limitations. In parallel, we intend to test the tool on additional datasets to further evaluate its effectiveness and efficiency.

To date, the base version of the tool has been tested on a dataset from the Agricultural Structure Survey on land use and livestock in Germany for 2020. It would be valuable to test the new version of the tool on data from before 2020 to assess whether the improvements we have made enhance the tool's effectiveness.

We plan to utilize the timestamps documented in the output CSV files to compare them with new timestamps obtained from re-testing. This comparison will help us determine whether the student assistants can work faster with the new version of the tool.

6 CONCLUSION

In this paper, we introduced the Spreadsheet Data Extractor (SDE) [?], an enhanced tool that builds upon the foundational work of Alexander Aue et al. [1]. By addressing key limitations of the existing solution, we implemented significant performance optimizations and usability enhancements. Specifically, SDE employs incremental loading of worksheets and optimizes rendering by processing only the visible cells, resulting in performance improvements that enable the tool to open large Excel files.

We also integrated the selection hierarchy, worksheet view, and output preview into a unified interface, streamlining the data extraction process.

By adopting a user-centric approach that gives users full control over data selection and metadata hierarchy definition without requiring programming knowledge, we provide a robust and accessible solution for data extraction. Our tool offers user-friendly features such as the ability to duplicate hierarchies of columns and tables and to move them over similar structures for reuse, reducing the need for repetitive configurations.

By contributing our improved Spreadsheet Data Extractor under the GNU General Public License v3.0, we enable the community to access, use, and enhance the tool freely, fostering collaboration and further innovation. By combining the strengths of the original approach with our enhancements in user interface and performance optimizations, our tool significantly improves the efficiency and reliability of data extraction from diverse and complex spreadsheet formats.

REFERENCES

- [1] Andrea Ackermann Alexander Aue. 2024. Converting data organised for visual perception into machine-readable formats. In 44. GIL-Jahrestagung, *Biodiversität fördern durch digitale Landwirtschaft*. Gesellschaft für Informatik eV, 179–184.

- 736 [2] Daniel W Barowy, Sumit Gulwani, Ted Hart, and Benjamin Zorn. 2015. FlashRelate: extracting relational data from
737 semi-structured spreadsheets using examples. *ACM SIGPLAN Notices* 50, 6 (2015), 218–228.
- 738 [3] D.J. Berndt, J.W. Fisher, A.R. Hevner, and J. Studnicki. 2001. Healthcare data warehousing and quality assurance.
739 *Computer* 34, 12 (2001), 56–65. <https://doi.org/10.1109/2.970578>
- 740 [4] Zhe Chen, Michael Cafarella, Jun Chen, Daniel Prevo, and Junfeng Zhuang. 2013. Senbazuru: A prototype spreadsheet
741 database management system. *Proceedings of the VLDB Endowment* 6, 12, 1202–1205.
- 742 [5] Haoyu Dong, Shijie Liu, Shi Han, Zhouyu Fu, and Dongmei Zhang. 2019. Tablesense: Spreadsheet table detection with
743 convolutional neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 33. 69–76.
- 744 [6] Angus Dunn. 2010. Spreadsheets-the Good, the Bad and the Downright Ugly. *arXiv preprint arXiv:1009.5705* (2010).
- 745 [7] Julian Eberius, Christoper Werner, Maik Thiele, Katrin Braunschweig, Lars Dannecker, and Wolfgang Lehner. 2013.
746 DeExelerator: a framework for extracting relational data from partially structured documents. In *Proceedings of the
747 22nd ACM international conference on Information & Knowledge Management*. 2477–2480.
- 748 [8] Elvis Koci, Dana Kuban, Nico Luettig, Dominik Olwig, Maik Thiele, Julius Gonsior, Wolfgang Lehner, and Oscar
749 Romero. 2019. Xlindy: Interactive recognition and information extraction in spreadsheets. In *Proceedings of the ACM
750 Symposium on Document Engineering 2019*. 1–4.
- 751 [9] Kate Lovett. 2023. two_dimensional_scrollables package - Commit 4c16f3e. <https://github.com/flutter/packages/commit/4c16f3ef40333aa0aebe8a1e46ef7b9fef9a1c1f> Accessed: 2023-08-17.
- 752 [10] Gursharan Singh, Leah Findlater, Kentaro Toyama, Scott Helmer, Rikin Gandhi, and Ravin Balakrishnan. 2009. Nu-
753 metric paper forms for NGOs. In *2009 International Conference on Information and Communication Technologies and
754 Development (ICTD)*. IEEE, 406–416.
- 755 [11] Statistisches Bundesamt (Destatis). 2024. *Statistischer Bericht - Pflegekraeftevorausberechnung - 2024 bis 2070*. Technical
756 Report. Statistisches Bundesamt, Wiesbaden, Germany. <https://www.destatis.de/DE/Themen/Gesellschaft-Umwelt/Bevoelkerung/Bevoelkerungsvorausberechnung/Publikationen/Downloads-Vorausberechnung/statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.html> Statistical Report - Projection of Nursing Staff - 2024 to 2070.
- 757 [12] Statistisches Bundesamt (Destatis). 2024. Statistischer Bericht: Rechnungsergebnis der Kernhaushalte der Gemeinden.
758 <https://www.destatis.de/DE/Themen/Staat/Oeffentliche-Finanzen/Ausgaben-Einnahmen/Publikationen/Downloads-Ausgaben-und-Einnahmen/statistischer-bericht-rechnungsergebnis-kernhaushalt-gemeinden-2140331217005.html> Accessed: 2024-11-29.
- 759 [13] Hannah West and Gina Green. 2008. Because excel will mind me! the state of constituent data management in small
760 nonprofit organizations. In *Proceedings of the Fourteenth Americas Conference on Information Systems*. Association for
761 Information Systems, AIS Electronic Library (AISel). <https://aisel.aisnet.org/amcis2008/336>

763 Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009
764

765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784