

1 **Spreadsheet Data Extractor (SDE): A Performance-Optimized,**
2 **User-Centric Tool for Transforming Semi-Structured Excel**
3 **Spreadsheets into Relational Data**

4
5 ANONYMOUS AUTHOR(S)
6
7 SUBMISSION ID:

8 Spreadsheets are ubiquitous tools used across various domains. Despite their widespread use, analyzing and
9 utilizing data stored in spreadsheets poses significant challenges due to their semi-structured nature. Data
10 in spreadsheets are often formatted primarily for human readability, employing layouts and styles that are
11 easily understood by people but are difficult for automated systems to interpret. While these unstructured
12 formats offer advantages—such as providing an easily comprehensible hierarchy of metadata—they complicate
13 automated data extraction. This paper introduces the Spreadsheet Data Extractor (SDE), an open-source tool
14 designed to convert semi-structured spreadsheet data into structured formats without requiring programming
15 knowledge. Building upon previous work, we have enhanced the SDE with incremental loading of worksheets,
16 accurate rendering of cell dimensions by directly parsing the Excel file's XML content, and performance
17 optimizations to handle large datasets efficiently. We compare our tool with existing solutions and demonstrate
18 its effectiveness through performance evaluations, highlighting its potential to facilitate efficient and reliable
19 data extraction from diverse spreadsheet formats.

20 CCS Concepts: • **Applied computing** → Spreadsheets; • **Information systems** → Data cleaning; •
21 **Software and its engineering** → Extensible Markup Language (XML); • **Human-centered computing**
22 → Graphical user interfaces; • **Theory of computation** → Data compression.

23 Additional Key Words and Phrases: Spreadsheets, Data cleaning, Relational Data, Excel, XML, Graphical user
24 interfaces
25

26 **ACM Reference Format:**

27 Anonymous Author(s). 2025. Spreadsheet Data Extractor (SDE): A Performance-Optimized, User-Centric Tool
28 for Transforming Semi-Structured Excel Spreadsheets into Relational Data. In *Proceedings of 2025 International*
29 *Conference on Management of Data (ACM PODS '25)*. ACM, New York, NY, USA, 17 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

31 **1 INTRODUCTION**

32 Spreadsheets are ubiquitous tools used across various domains, including healthcare [5], nonprofit
33 organizations [16, 21], finance, commerce, academia, and government [10]. Despite their widespread
34 use, analyzing and utilizing data stored in spreadsheets poses significant challenges for data
35 analysis and utilization, primarily due to their semi-structured nature. Data within spreadsheets
36 are frequently organized for human readability [20], featuring layouts with empty cells, merged
37 cells, hierarchical headers, and multiple tables. While these formats enhance comprehensibility for
38 users [15], they impede machine readability and automated data processing. This semi-structured
39

40
41 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee
42 provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the
43 full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored.
44 Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires
45 prior specific permission and/or a fee. Request permissions from permissions@acm.org.

46 *ACM PODS '25, June 22–27, 2025, Berlin, Germany*

47 © 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

48 ACM ISBN ...\$15.00

49 <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

organization complicates the extraction of meaningful insights, especially when attempting to integrate spreadsheet data into more robust and scalable data management systems.

The reliance on spreadsheets as ad-hoc solutions poses several limitations:

- **Data Integrity and Consistency:** Spreadsheets are prone to errors, such as duplicate entries, inconsistent data formats, and inadvertent modifications, which can compromise data integrity. [1, 7]
- **Scalability Issues:** As datasets grow in size and complexity, spreadsheets become less efficient for data storage and retrieval, leading to performance bottlenecks. [14, 19]
- **Limited Query Capabilities:** Unlike databases, spreadsheets lack advanced querying and indexing features, restricting users from performing complex data analyses.

Transitioning from these ad-hoc spreadsheet solutions to standardized database systems offers numerous benefits:

- **Enhanced Data Integrity:** Databases enforce data validation rules and constraints, ensuring higher data quality and consistency.
- **Improved Scalability:** Databases are designed to handle large volumes of data efficiently, supporting complex queries and transactions without significant performance degradation.
- **Advanced Querying and Reporting:** Databases provide powerful querying languages like SQL, enabling sophisticated data analysis and reporting capabilities.
- **Seamless Integration:** Databases facilitate easier integration with various applications and services, promoting interoperability and data sharing across platforms.

Given the abundance of existing spreadsheet data and the clear advantages of database systems, there is a pressing need for tools that can bridge the gap between these two formats. Automated and accurate data extraction from spreadsheets into relational database formats is essential for organizations to leverage their data assets effectively.

Previous work by Aue et al. introduced a tool that facilitates data extraction from Excel files [3]. While effective, their solution faced performance issues and inaccuracies in rendering cell dimensions, limiting its usability with large and complex datasets.

In this paper, we aim to transform semi-structured spreadsheet data into machine-readable formats by building upon their work. We present the Spreadsheet Data Extractor (SDE), an enhanced tool that enables users to define data hierarchies through cell selection without any programming knowledge. [2] Our enhancements address key limitations of the existing solution, making data extraction from complex spreadsheets more efficient and user-friendly.

1.1 Contributions

Our main contributions are as follows:

- (1) We release the SDE under the open-source GNU General Public License v3.0, promoting community access and collaboration [2].
- (2) We implement incremental loading of worksheets to enhance performance, allowing the tool to handle large Excel files efficiently.
- (3) We accurately render row heights and column widths by parsing XML data, ensuring that the spreadsheet's visual representation closely matches that of Excel.
- (4) We optimize the rendering engine to draw only the visible cells, significantly improving performance when dealing with large datasets.
- (5) We integrate the selection hierarchy, worksheet view, and output preview into a unified interface, streamlining the user experience.

99 2 RELATED WORK

100 The extraction of relational data from semi-structured documents, particularly spreadsheets, has
101 garnered significant attention due to their ubiquitous use across domains such as business, gov-
102 ernment, and scientific research. Several frameworks and tools have been developed to address
103 the challenges of converting flexible spreadsheet formats into normalized relational forms suitable
104 for data analysis and integration. Notable among these are **DeExcelerator** [11], **XLIndy** [12],
105 **FLASHRELATE** [4], **Senbazuru** [6], **TableSense** [9], and the approach by Aue et al. [3], on which
106 our work builds.

107

108 2.1 Aue et al.'s Converter

109 Aue et al. [3] developed a tool to facilitate data extraction from Excel spreadsheets by leveraging
110 the Dart excel package [8] to process .xlsx files. This tool allows users to define data hierarchies
111 by selecting relevant cells containing data and metadata. However, the approach faced significant
112 performance bottlenecks due to the excel package's requirement to load the entire .xlsx file into
113 memory, resulting in slow response times, particularly for large files.

114 In addition to memory issues, the tool calculated row heights and column widths based solely
115 on cell content, ignoring the dimensions specified in the original Excel file. This led to rendering
116 discrepancies between the tool and the original spreadsheet. Furthermore, the tool rendered all
117 cells, regardless of their visibility within the viewport, significantly degrading performance when
118 handling worksheets with large numbers of cells.

119

120 2.2 DeExcelerator

121 Eberius et al. [11] introduced **DeExcelerator**, a framework that transforms partially structured
122 spreadsheets into first normal form relational tables using heuristic-based extraction phases. It
123 addresses challenges such as table detection, metadata extraction, and layout normalization. While
124 effective in automating normalization, its reliance on predefined heuristics limits adaptability to
125 heterogeneous or unconventional spreadsheet formats, highlighting the need for more flexible
126 approaches.

127

128 2.3 XLIndy

129 Koci et al. [12] developed **XLIndy**, an interactive Excel add-in with a Python-based machine
130 learning backend. Unlike DeExcelerator's fully automated heuristic approach, XLIndy integrates
131 machine learning techniques for layout inference and table recognition, enabling a more adaptable
132 and accurate extraction process. XLIndy's interactive interface allows users to visually inspect
133 extraction results, adjust configurations, and compare different extraction runs, facilitating iterative
134 fine-tuning. Additionally, users can manually revise predicted layouts and tables, saving these
135 revisions as annotations to improve classifier performance through (re-)training. This user-centric
136 approach enhances the tool's flexibility, allowing it to accommodate diverse spreadsheet formats
137 and user-specific requirements more effectively than purely heuristic-based systems.

138

139 2.4 FLASHRELATE

140 Barowy et al. [4] presented **FLASHRELATE**, an approach that empowers users to extract struc-
141 tured relational data from semi-structured spreadsheets without requiring programming expertise.
142 FLASHRELATE introduces a domain-specific language, **FLARE**, which extends traditional regular
143 expressions with spatial constraints to capture the geometric relationships inherent in spreadsheet
144 layouts. Additionally, FLASHRELATE employs an algorithm that synthesizes FLARE programs

from a small number of user-provided positive and negative examples, significantly simplifying the automated data extraction process.

FLASHRELATE distinguishes itself from both DeExcelerator and XLIIndy by leveraging programming-by-example (PBE) techniques. While DeExcelerator relies on predefined heuristic rules and XLIIndy incorporates machine learning models requiring user interaction for fine-tuning, FLASHRELATE allows non-expert users to define extraction patterns through intuitive examples. This approach lowers the barrier to entry for extracting relational data from complex spreadsheet encodings, making the tool accessible to a broader range of users.

2.5 Senbazuru

Chen et al. [6] introduced **Senbazuru**, a prototype Spreadsheet Database Management System (SSDBMS) designed to extract relational information from a large corpus of spreadsheets. Senbazuru addresses the critical issue of integrating data across multiple spreadsheets, which often lack explicit relational metadata, thereby hindering the use of traditional relational tools for data integration and analysis.

Senbazuru comprises three primary functional components:

(1) **Search**: Utilizing a textual search-and-rank interface, Senbazuru enables users to quickly locate relevant spreadsheets within a vast corpus. The search component indexes spreadsheets using Apache Lucene, allowing for efficient retrieval based on relevance to user queries.

(2) **Extract**: The extraction pipeline in Senbazuru consists of several stages:

- **Frame Finder**: Identifies data frame structures within spreadsheets using Conditional Random Fields (CRFs) to assign semantic labels to non-empty rows, effectively detecting rectangular value regions and associated attribute regions.

- **Hierarchy Extractor**: Recovers attribute hierarchies for both left and top attribute regions. This stage also incorporates a user-interactive repair interface, allowing users to manually correct extraction errors, which the system then generalizes to similar instances using probabilistic methods.

- **Tuple Builder and Relation Constructor**: Generates relational tuples from the extracted data frames and assembles these tuples into coherent relational tables by clustering attributes and recovering column labels using external schema repositories like Freebase and YAGO.

(3) **Query**: Supports basic relational operations such as selection and join on the extracted relational tables, enabling users to perform complex data analysis tasks without needing to write SQL queries.

Senbazuru's ability to handle hierarchical spreadsheets, where attributes may span multiple rows or columns without explicit labeling, sets it apart from earlier systems like DeExcelerator and XLIIndy. By employing machine learning techniques and providing user-friendly repair interfaces, Senbazuru ensures high-quality extraction and facilitates the integration of spreadsheet data into relational databases.

2.6 TableSense

Dong et al. [9] developed **TableSense**, an end-to-end framework for spreadsheet table detection using Convolutional Neural Networks (CNNs). TableSense addresses the diversity of table structures and layouts by introducing a comprehensive cell featurization scheme, a Precise Bounding Box Regression (PBR) module for accurate boundary detection, and an active learning framework to efficiently build a robust training dataset.

197 While **DeExcelerator**, **XLIndy**, **FLASHRELATE**, and **Senbazuru** focus primarily on transforming
198 spreadsheet data into relational forms through heuristic, machine learning, and programming-
199 by-example approaches, **TableSense** specifically targets the accurate detection of table boundaries
200 within spreadsheets using deep learning techniques. Unlike region-growth-based methods em-
201 ployed in commodity spreadsheet tools, which often fail on complex table layouts, TableSense
202 achieves superior precision and recall by leveraging CNNs tailored for the unique characteristics of
203 spreadsheet data. However, TableSense focuses on table detection and visualization, allowing users
204 to generate diagrams from the detected tables but does not provide functionality for exporting the
205 extracted data for further analysis.

207 2.7 Comparison and Positioning

208 While **DeExcelerator**, **XLIndy**, **FLASHRELATE**, **Senbazuru**, and **TableSense** each offer unique
209 approaches to spreadsheet data extraction, they share certain limitations. Many of these tools
210 are not readily accessible: **FLASHRELATE** and **TableSense** are proprietary, and **Senbazuru**,
211 **XLIndy**, and **DeExcelerator** are discontinued projects with limited or no source code availability.
212 In contrast, we contribute our spreadsheet data extractor under the GNU General Public License
213 v3.0, allowing the community to access, use, and improve the tool freely.

214 Moreover, unlike the aforementioned tools that rely on heuristics, machine learning, or AI
215 techniques—which can introduce errors requiring users to identify and correct—we adopt a user-
216 centric approach that gives users full control over data selection and metadata hierarchy definition.
217 While this requires more manual input, it eliminates the uncertainty and potential inaccuracies
218 associated with automated methods. To streamline the process and enhance efficiency, our tool
219 includes user-friendly features such as the ability to duplicate hierarchies of columns and tables,
220 and to move them over similar structures for reuse, reducing the need for repetitive configurations.
221

222 By combining the strengths of manual control with enhanced user interface features and per-
223 formance optimizations, our tool offers a robust and accessible solution for extracting relational
224 data from complex and visually intricate spreadsheets. These enhancements not only improve
225 performance and accuracy but also elevate the overall user experience, making our tool a valuable
226 asset for efficient and reliable data extraction from diverse spreadsheet formats.

227 3 METHODOLOGY

228 In this section, we detail the design and implementation of the Spreadsheet Data Extractor (SDE),
229 emphasizing its user-centric approach and performance optimizations. The SDE enables users to
230 transform semi-structured spreadsheet data into structured, machine-readable formats without
231 requiring programming expertise. We achieve this through an intuitive interface that allows for
232 cell selection and hierarchy definition, incremental loading of worksheets, accurate rendering of
233 cell dimensions, and optimized performance for handling large datasets by incrementally loading
234 worksheets and by rendering only the cells that are currently visible in the view.
235

236 3.1 User-Centric Data Extraction

237 The core functionality of the SDE revolves around allowing users to select cells containing data
238 and metadata to define a data hierarchy. This process is facilitated through a graphical interface
239 that displays the spreadsheet and allows for intuitive selection and manipulation of the selection
240 hierarchy.

241 242 *3.1.1 Hierarchy Definition.* Users can select individual cells or ranges of cells by clicking and using
243 shift-click for multi-selection. These selections represent either data or metadata.

The selected cells are organized into a hierarchical tree structure, where each node represents a data element, and child nodes represent nested data or metadata. This hierarchy defines how the data will be transformed into a structured format.

3.1.2 Reusability and Efficiency. To optimize the extraction process and reduce repetitive tasks, the SDE allows users to duplicate previously defined hierarchies and apply them to similar regions within the spreadsheet. This feature is particularly useful for spreadsheets with repeating structures, such as multiple tables with the same format.

3.2 Example Workflow

Consider a spreadsheet containing statistical forecasts of future nursing staff availability in Germany [17]. Figure 1 shows the SDE interface, which consists of three main components:

Hierarchy Panel (Top Left): Displays the hierarchy of cell selections, initially empty.

Spreadsheet View (Top Right): Shows the currently opened Excel file and the currently selected worksheet for cell selection.

Output Preview (Bottom): Provides immediate feedback on the data extraction based on current selections.

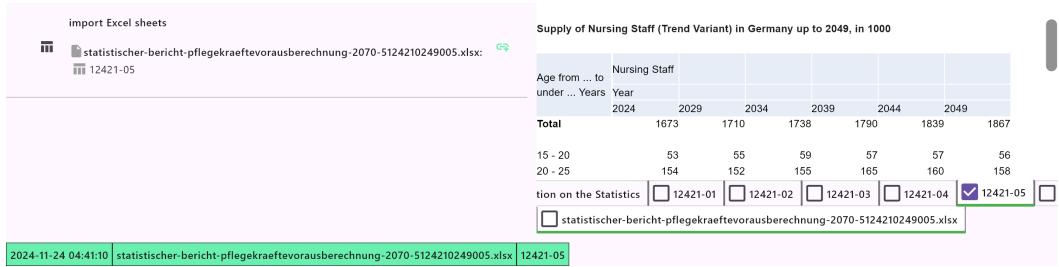


Fig. 1. The SDE Interface Overview.

3.2.1 Selection of the First Column. The user adds a node to the hierarchy and selects the cell containing the metadata "Nursing Staff" (Figure 2). This cell represents metadata that is common to all cells in this worksheet. Therefore, it should be selected first and should appear at the beginning of each row in the output CSV file.

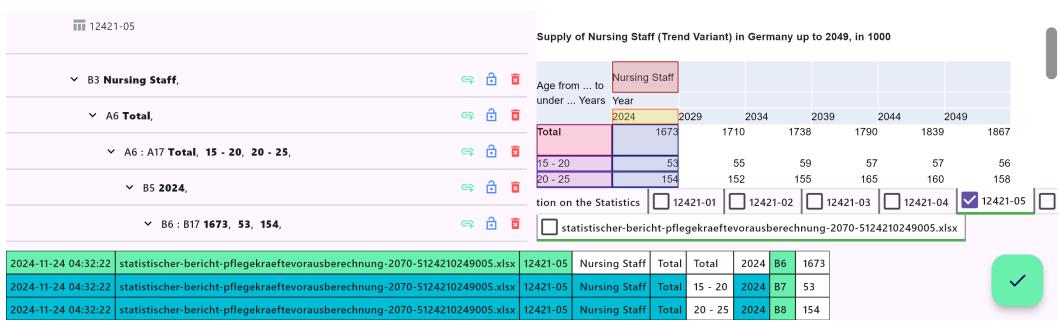


Fig. 2. Selection of the First Column Metadata

Within this node, the user adds a child node and selects the cell "*Total*", which serves as both a table header and a row label. This selection represents the table header of the first subtable. The user

295 adds another child node and selects the range of cells containing row labels (e.g., "Total", "15-20",
 296 "20-25" and so forth) by clicking the first cell and shift-clicking the last cell.

297 A further child node is then placed under the row labels node, and the user selects the year
 298 "2024". Subsequently, an additional child node is created beneath the year node, and the user selects
 299 the corresponding data cells (e.g., "1673", "53", "154", etc.).

300 At this point, the hierarchy consists of five nodes, each—except the last one—containing an
 301 embedded child node. In the upper-right portion of the interface, the chosen cells are displayed
 302 in distinct colors corresponding to each node. The lower area shows a preview of the extracted
 303 output. For each child node, an additional column is appended to the output. When multiple cells
 304 are selected for a given node, their values appear as entries in new rows of the output, reflecting
 305 the defined hierarchical structure.

306 3.2.2 *Duplicating the Column Hierarchy.* To avoid repetitive manual entry for additional years, the
 307 user duplicates the hierarchy for "2024" and adjusts the cell selections to include data for subsequent
 308 years (e.g., "2025", "2026") using the "Move and Duplicate" feature.

309 310 To do this, the user selects the node of the first column "2024" and right-clicks on it. A popup
 311 opens in which the action "move and duplicate" appears, which should then be clicked, as shown
 312 in Figure 3a.

The figure consists of two screenshots of a spreadsheet application interface.
 (a) On the left, a context menu is open over a cell in the 'Year' column. The menu items include 'encapsulate', 'enter value manually', 'move and duplicate' (which is checked), 'copy cell selection', 'copy task', and 'duplicate task'. The 'move and duplicate' option is highlighted.
 (b) On the right, a dialog box titled 'Supply of Nursing Staff (Trend Variant) in Germany up to 2049, in 1000' is shown. It contains a table with data for 'Nursing Staff' across years 2024 to 2049. Below the table, there are checkboxes for 'move and duplicate' and a 'repeat' input field set to '5'. There are also buttons for 'cancel' and 'ok'. The bottom part of the dialog shows a preview of the replicated data, with the first few rows matching the original table and subsequent rows showing the shifted and duplicated data.

323 (a) Invoking the "move and duplicate" feature on the
 324 325 2024 column node. (b) Adjusting the number of repetitions to duplicate
 326 the column selection.

327 Fig. 3. Utilizing the "Move and Duplicate" feature to replicate column hierarchies for additional years.

328 329 Subsequently, a series of buttons opens in the app bar at the top right, allowing the user to move
 330 the cell selections of the node as well as all child nodes, as seen in Figure 3b. By pressing the button
 331 to move the selection by one unit to the right, the next column is selected; however, this would
 332 also deselect the first column since the selection was moved. To preserve the first column, the
 333 "move and duplicate" checkbox can be activated. This creates the shifted selection in addition to the
 334 original selection. However, the changes are only applied when the accept button is clicked. The
 335 next columns could also be selected in the same way. But this can be done faster, because instead
 336 of moving the selection and duplicating it only once, the "repeat" input field can be filled with as
 337 many repetitions as there are columns. By entering the number 5, the selection of the first column
 338 is shifted 5 times by one unit to the right and duplicated at each step.

339 340 The user reviews the selections in the spreadsheet view, where each selection is highlighted in a
 341 different color corresponding to its node in the hierarchy. Only after the user has reviewed the
 342 343

shifted and duplicated selections in the worksheet and clicked the accept button are the nodes in the hierarchy created as desired. We present the results of these runtime measurements in Figure 4.

The user reviews the selections in the spreadsheet view, where each selection is highlighted in a different color corresponding to its node in the hierarchy. Erst nachdem der Nutzer die verschobenen und duplizierten selektionen in der Worksheet ansieht überprüft hat und den Akzeptieren button geklickt hat, werden die Knoten in der Hierachie wie gewünscht angelegt. Das Ergebnis dieser Operation ist in Abbildung 4 zu sehen.

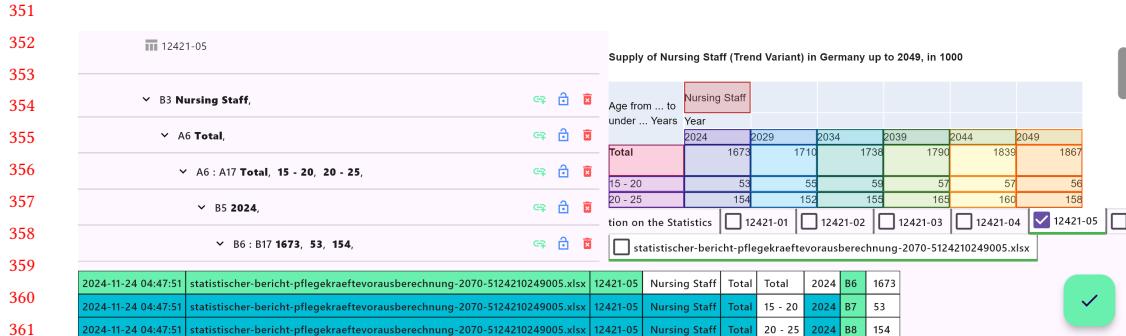


Fig. 4. Resulting Hierarchy After Move and Accept

3.2.3 Duplicating the Table Hierarchy. The same method that worked effectively for duplicating the columns can now be applied to the subtables, as shown in Figure 5.



Fig. 5. Selection of All Cells in the Subtables by Duplicating the Hierarchy of the First Table

By selecting the node with the value "Total" and clicking the "Move and Duplicate" button, we can apply the selection of the "Total" subtable to the other subtables. This involves shifting the table downward by as many rows as necessary to overlap with the subtable below.

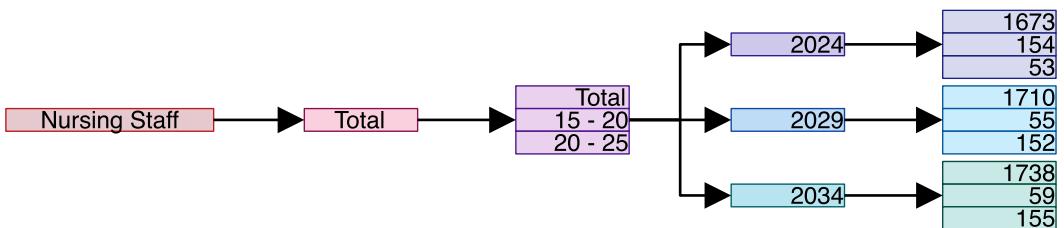
393 However, there is a minor issue: the child nodes of the "Total" node also include the column
 394 headers. If these column headers were repeated in the subtables below, shifting the selections
 395 downward would work without modification. Since these cells are not repeated in the subtables,
 396 we need to prevent the column headers cells from moving during the duplication process.

397 To achieve this, we can exclude individual nodes from being moved by locking their selection.
 398 This is done by clicking the padlock icon on the corresponding nodes, which freezes their cell
 399 selection and keeps them fixed at their original position, regardless of other cells being moved.

400 Therefore, we identify and select the nodes containing the column headers—specifically, the
 401 years 2024 to 2049—and lock their selection using the padlock button. By shifting the selection
 402 downward and duplicating it, we can easily move and duplicate the cell selections for the subtables
 403 below. By setting the number of repetitions to 2, all subtables are completely selected.

404 405 3.3 Cross Product Transformation

406 The graph resulting from the selected hierarchy is shown in Figure 6. To simplify the explanation,
 407 the example is limited to the first three columns and the first three rows of the first sub-table.



417 Fig. 6. Illustration of the Cross Product Transformation Before Application

418 Once the hierarchy is defined, the **Structured Data Engine (SDE)** applies a cross-product
 419 transformation to generate a relational format from the selection graph. This transformation
 420 consists of two key steps:

- 421 (1) **Node Duplication:** Nodes with multiple incoming or outgoing edges (e.g., the row labels
 422 node with the values *Total*, *15-20*, *20-25*) are duplicated to ensure that each edge connects to a
 423 unique instance of the node. This adjustment replaces the original many-to-one relationships
 424 with one-to-one mappings for each edge.
- 425 (2) **Value Replication:** Nodes containing single values (e.g., the year *2024*) are replicated
 426 to align with the number of values associated with the connected node. This ensures a
 427 consistent structure in the relational output, maintaining alignment across all hierarchical
 428 levels.

429 The resulting graph after the cross-product transformation is shown in Figure 7. With this
 430 transformation applied to the selected hierarchy, the SDE generates a structured output that reflects
 431 the original spreadsheet's hierarchical relationships, enabling users to analyze and integrate the
 432 data effectively.

433 434 3.4 Incremental Loading of Worksheets

435 Opening large Excel files traditionally involves loading the entire file and all its worksheets into
 436 memory before displaying any content. In files containing very large worksheets, this process
 437 can take several seconds to minutes, causing significant delays for users who need to access data
 438 quickly.

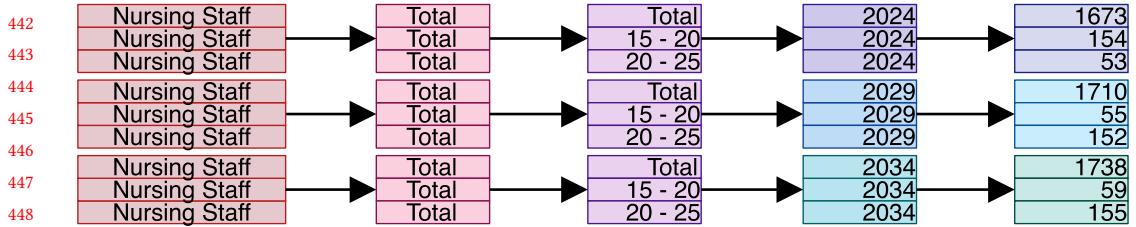


Fig. 7. Illustration of the Cross Product Transformation After Application

To facilitate efficient data extraction from multiple Excel files, we implemented a mechanism for incremental loading of worksheets within the SDE. Excel files (.xlsx format) are ZIP archives containing a collection of XML files that describe the worksheets, styles, and shared strings. Key components include:

- **xl/sharedStrings.xml**: Contains all the unique strings used across worksheets, reducing redundancy.
- **xl/styles.xml**: Defines the formatting styles for cells, including fonts, colors, and borders.
- **xl/worksheets/sheetX.xml**: Represents individual worksheets (sheet1.xml, sheet2.xml, etc.).

Our solution opens the Excel file as a ZIP archive and initially extracts only the essential metadata and shared resources required for the application to function. This initial extraction includes:

(1) **Metadata Extraction:**

We read the archive's directory to identify the contained files without decompressing them fully. This step is quick, taking only a few milliseconds, and provides information about the available worksheets and shared resources.

(2) **Selective Extraction:**

We immediately extract `sharedStrings.xml` and `styles.xml` because these files are small and contain information necessary for rendering cell content and styles across all worksheets. These files are parsed and stored in memory for quick access during rendering.

(3) **Deferred Worksheet Loading:**

The individual worksheet files (`sheetX.xml`) remain compressed and are loaded into memory in their binary unextracted form. They are not decompressed or parsed at this stage.

(4) **On-Demand Parsing:**

When a user accesses a specific worksheet—either by selecting it in the interface or when a unit test requires data from it—the corresponding `sheetX.xml` file is then decompressed and parsed. This parsing occurs in the background and is triggered only by direct user action or programmatic access to the worksheet's data.

(5) **Memory Release:**

After a worksheet has been decompressed and its XML parsed, we release the memory resources associated with the parsed data. This approach prevents excessive memory usage and ensures that the application remains responsive even when working with multiple large worksheets.

By adopting this incremental loading approach, users experience minimal wait times when opening an Excel file. The initial loading is nearly instantaneous, allowing users to begin interacting with the application without delay. This contrasts with traditional methods that require loading all worksheets upfront, leading to significant wait times for large files.

491 **3.5 Rendering of Worksheets**

492 To ensure that users can navigate worksheets without difficulty, we prioritize displaying the
 493 worksheets in a manner that closely resembles their appearance in Excel. This involves accurately
 494 rendering cell dimensions, formatting, and text behaviors.

495 **3.5.1 Displaying Row Heights and Column Widths.** Our solution extracts information about column
 496 widths and row heights directly from the Excel file's XML structure. Specifically, we retrieve the
 497 column widths from the *width* attribute of the `<col>` elements and the row heights from the *ht*
 498 attribute of the `<row>` elements in the `sheetX.xml` files.
 499

500 In Excel, column widths and row heights use units that do not directly map to pixels, requiring
 501 conversion for accurate on-screen rendering. Different scaling factors are needed for columns and
 502 rows. Through empirical testing, we derived the following scaling factors:

- **Column Widths:** Multiply the *width* attribute by 7.
- **Row Heights:** Multiply the *ht* attribute by $\frac{4}{3}$.

503 Despite research, we could not find official documentation explaining the rationale behind these
 504 specific scaling factors. This lack of documentation poses a challenge for accurately replicating
 505 Excel's rendering.

506 **3.5.2 Cell Formatting.** Cell formatting plays a crucial role in accurately representing the appearance
 507 of worksheets. Formatting information is stored in the `styles.xml` file, where styles are defined
 508 and later referenced in the `sheetX.xml` files as shown in Figure 8.

```

512
513   <font count="13">
514     <!-- ... -->
515     <font>
516       <sz val="10"/>
517       <name val="Arial"/>
518       <family val="2"/>
519     </font>
520     <!-- ... -->
521   </fonts>
522
523   <cellXfs count="19">
524     <!-- ... -->
525     <xf numFmtId="0" fontId="2" fillId="2" borderId="1" xfId="1" applyFill="1" applyBorder="1" applyAlignment="1">
526       <alignment horizontal="centerContinuous" vertical="center" wrapText="1"/>
527     </xf>
528   </cellXfs>
529
530   xl/styles.xml
531
532
533   <row r="3" spans="1:7" s="2" customFormat="1" ht="25.5" customHeight="1">
534     <c r="B3" s="10" t="s"><v>17</v></c>
535   </row>
536
537   xl/worksheets/sheetX.xml
538
539

```

Fig. 8

537 Each cell in the worksheet references a style index through the *s* attribute, which points to the
 538 corresponding `<xf>` element within the `cellXfs` collection. These `<xf>` elements contain attributes

such as `fontId`, `fillId`, and `borderId`, which reference specific font, fill (background), and border definitions located in the `fonts`, `fills`, and `borders` collections, respectively. By parsing these references, we can accurately apply the appropriate fonts, background colors, and border styles to each cell.

Through meticulous parsing and application of these formatting details, we ensure that the rendered worksheet closely mirrors the original Excel file, preserving the visual cues and aesthetics that users expect.

3.5.3 Handling Text Overflow. In Excel, when the content of a cell exceeds its width, the text may overflow into adjacent empty cells, provided those cells do not contain any data. If adjacent cells are occupied, Excel truncates the overflowing text at the cell boundary. Replicating this behavior is essential for accurate rendering and user familiarity.

We implemented text overflow handling by checking if the adjacent cell to the right is empty before allowing text to overflow. If the adjacent cell is empty, we extend the text rendering. If the adjacent cell contains data, we truncate the text at the boundary of the original cell.

Figure 1 illustrates this behavior. The text "Supply of Nursing Staff ..." extends into the neighboring cell because it is empty. If not for this handling, the text would be truncated at the cell boundary, leading to incomplete data display as shown in Figure 9.

[To the table of contents](#)

Supply of Nurs

Age from ... to under ... Years	Nursing Staff	2024	2029	2034	2039	2044	2049	
Total		1673	1710	1738	1790	1839	1867	

Fig. 9. Falsche Darstellung ohne overflow von zellen mit angrenzenden zellen ohne inhalt

By accurately handling text overflow, we improve readability and maintain consistency with Excel's user interface, which is crucial for users transitioning between Excel and our tool.

3.6 Performance Optimization

To ensure high frame rates even with large worksheets, we optimized the Spreadsheet Data Extractor to render only the cells that are currently visible to the user. Rendering the entire worksheet, especially when it contains thousands of cells, can significantly degrade performance. By focusing on the visible cells within the viewport, we reduce computational overhead and improve responsiveness.

We achieve this optimization by utilizing the `two_dimensional_scrollables` package [13]. This package provides functionality for efficiently handling two-dimensional scrolling regions, making it suitable for rendering large grids like spreadsheets.

Since we extract all column widths and row heights from the XML files, we can calculate the exact dimensions and positions of each cell. By accumulating the widths and heights, we determine the

589 coordinates of each cell within the worksheet grid. These coordinates are essential for identifying
590 which cells fall within the current viewport and should be rendered.

591 We consider the current scroll offsets along the horizontal (x -axis) and vertical (y -axis) directions.
592 The viewport is defined by the width and height of the panel displaying the Excel worksheet. To
593 determine the visible cells, we perform the following steps:

594 • **Horizontal Visibility:**

- 595 – Sum the column widths until the accumulated sum reaches the left edge of the viewport.
596 All cells to the left are ignored.
- 597 – Continue adding column widths until the sum exceeds the right edge of the viewport.
598 Cells beyond this point are also ignored.

599 • **Vertical Visibility:**

- 600 – Sum the row heights until the accumulated sum reaches the top edge of the viewport.
601 All cells above are ignored.
- 602 – Continue adding row heights until the sum surpasses the bottom edge of the viewport.
603 Cells below this point are ignored.

604 By rendering only the cells within these boundaries, we significantly reduce the number of cells
605 processed at any given time.

606 The `two_dimensional_scrollables` package provides interfaces where the logic for laying out
607 the cells can be implemented. Parameters that describe the viewport, including the horizontal and
608 vertical offsets as well as the viewport height and width, are supplied by these interfaces.

609 Algorithm 1 outlines our implementation of the overridden `layoutChildSequence` function,
610 which is invoked by the `two_dimensional_scrollables` package to calculate and arrange the
611 visible cells within the viewport.

612 By applying this method, we render only the cells necessary for the current view, thereby
613 optimizing performance and ensuring smooth user interactions even with large and complex
614 worksheets.

615 **4 EVALUATION**

616 The fundamental approach of the Spreadsheet Data Extractor, based on the converter by Aue et
617 al. [3], upon which we build, remains unchanged. The effectiveness of this approach has already
618 been investigated. Aue et al. evaluated the extraction of data from over 500 Excel files. The time
619 required for each file was determined from a sample of 331 processed Excel files comprising 3,093
620 worksheets. On average, student assistants needed 15 minutes per file and 95 seconds per worksheet.

621 Our focus is on improving the user experience and optimizing the performance of the Spreadsheet
622 Data Extractor. We enhanced the user experience by displaying the Excel worksheets similarly
623 to how they appear in Excel and by reducing the number of required user interactions through
624 the integration of the selection hierarchy, worksheet view, and output preview into a single
625 interface. Performance was further improved by implementing incremental loading of Excel files
626 and rendering only the visible cells.

627 **4.1 Acceleration When Opening Files**

628 To evaluate the performance improvements when opening Excel files, we conducted a series of tests
629 using a large Excel file. We downloaded the entire collection of Excel files from the German Federal
630 Statistical Office (Destatis)¹, which provides extensive statistical data across various domains. From
631 this collection, we identified the largest file [18], which has a compressed size of 87 MB and an
632 uncompressed size of 911 MB.

633 ¹<https://www.destatis.de>

Algorithm 1 Layout of Visible Spreadsheet Cells

```

638
639 1: Initialize Indices
640 2:   leadingColumnIndex  $\leftarrow$  column index corresponding to horizontalOffset
641 3:   leadingRowIndex  $\leftarrow$  row index corresponding to verticalOffset
642 4:   trailingColumnIndex  $\leftarrow$  column index corresponding to horizontalOffset + viewportWidth
643 5:   trailingRowIndex  $\leftarrow$  row index corresponding to verticalOffset + viewportHeight
644
645 6: Calculate Initial Offsets
646 7:   leadingColumnOffset  $\leftarrow$  sum of widths from the first column up to leadingColumnIndex
647 8:   leadingRowOffset  $\leftarrow$  sum of heights from the first row up to leadingRowIndex
648 9:   horizontalLayoutOffset  $\leftarrow$  leadingColumnOffset – horizontalOffset
649 10: for each columnIndex from leadingColumnIndex to trailingColumnIndex do
650 11:   verticalLayoutOffset  $\leftarrow$  leadingRowOffset – verticalOffset
651 12:   for each rowIndex from leadingRowIndex to trailingRowIndex do
652 13:     cell  $\leftarrow$  build or retrieve the cell at (columnIndex, rowIndex)
653 14:     Layout cell at position (horizontalLayoutOffset, verticalLayoutOffset)
654 15:     if a custom height is defined for row rowIndex then
655 16:       verticalLayoutOffset  $\leftarrow$  verticalLayoutOffset + height of row rowIndex
656 17:     else
657 18:       verticalLayoutOffset  $\leftarrow$  verticalLayoutOffset + defaultRowHeight
658 19:     end if
659 20:   end for
660 21:   columnWidth  $\leftarrow$  width of column columnIndex
661 22:   horizontalLayoutOffset  $\leftarrow$  horizontalLayoutOffset + columnWidth
662 23: end for
663

```

664 We performed three sets of tests to compare the performance of different methods for opening
 665 and reading data from this Excel file:

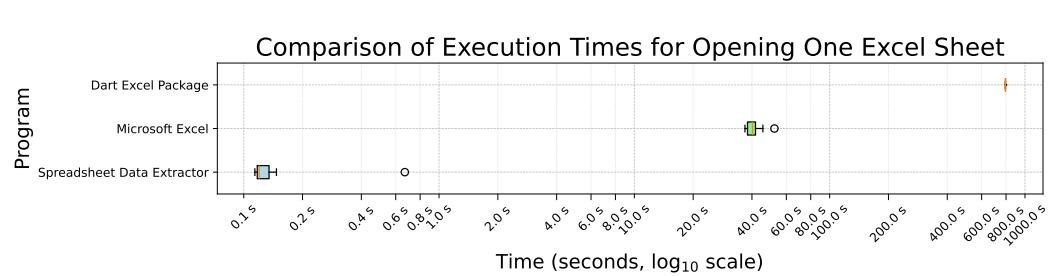
- 666 (1) **VBA Script in Excel:** We wrote a VBA script that opens the Excel file and reads specific
 667 values from a worksheet. This script simulates how users might interact with the file using
 668 Excel's built-in capabilities.
- 669 (2) **Spreadsheet Data Extractor:** We developed an equivalent unit test using the functions
 670 implemented in our Spreadsheet Data Extractor to open the same file and read the same
 671 cells. This test aimed to assess the performance of our tool under the same conditions.
- 672 (3) **excel Package:** We tested the excel package [8] used by Aue et al. [3], creating a unit test
 673 that attempts to read the same cells from the file. This test was intended to benchmark the
 674 performance of the prior solution.

676 All tests were conducted on a machine with the following specifications: Intel Core i5-10210U
 677 quad-core CPU at 1.60 GHz, 16 GB RAM, and a solid-state drive (SSD), running Windows 10 Pro.
 678 The version of Microsoft Excel used was Microsoft Office LTSC Professional Plus 2021.

679 We ran the script 10 times, measuring the time required to open the file and read the values
 680 in each iteration. The results of these runtime measurements are presented in Figure 10. The
 681 Spreadsheet Data Extractor opened the worksheet with a median time of 120 milliseconds and an
 682 average time of 178 milliseconds. The first run took 668 milliseconds, possibly because the file was
 683 not yet cached and had to be loaded from the disk. In contrast, Excel opened the worksheet with a
 684 median time of 40.281 seconds and an average time of 41.138 seconds. The Dart excel package [8]
 685 used in the previous work took 13 minutes and 15 seconds to open the worksheet in the first run.

687 The other nine runs could not be completed because an out-of-memory exception was thrown
 688 during the second run.

689 These results demonstrate that the Spreadsheet Data Extractor opens worksheets over two orders
 690 of magnitude faster than Excel and nearly four orders of magnitude faster than the Dart excel
 691 package used in prior work.



706 5 FUTURE WORK

707 We plan to continue improving the Spreadsheet Data Extractor by implementing new features that
 708 enhance the user experience and address current limitations. In parallel, we intend to test the tool
 709 on additional datasets to further evaluate its effectiveness and efficiency.

710 To date, the base version of the tool has been tested on a dataset from the Agricultural Structure
 711 Survey on land use and livestock in Germany for 2020. It would be valuable to test the new version
 712 of the tool on data from before 2020 to assess whether the improvements we have made enhance
 713 the tool's effectiveness.

714 We plan to utilize the timestamps documented in the output CSV files to compare them with
 715 new timestamps obtained from re-testing. This comparison will help us determine whether the
 716 student assistants can work faster with the new version of the tool.

717 6 CONCLUSION

718 In this paper, we introduced the Spreadsheet Data Extractor (SDE) [2], an enhanced tool that builds
 719 upon the foundational work of Alexander Aue et al. [3]. By addressing key limitations of the existing
 720 solution, we implemented significant performance optimizations and usability enhancements.
 721 Specifically, SDE employs incremental loading of worksheets and optimizes rendering by processing
 722 only the visible cells, resulting in performance improvements that enable the tool to open large
 723 Excel files.

724 We also integrated the selection hierarchy, worksheet view, and output preview into a unified
 725 interface, streamlining the data extraction process. By adopting a user-centric approach that gives
 726 users full control over data selection and metadata hierarchy definition without requiring program-
 727 ming knowledge, we provide a robust and accessible solution for data extraction. Our tool offers
 728 user-friendly features such as the ability to duplicate hierarchies of columns and tables and to move
 729 them over similar structures for reuse, reducing the need for repetitive configurations.

730 By combining the strengths of the original approach with our enhancements in user interface
 731 and performance optimizations, our tool significantly improves the efficiency and reliability of data
 732 extraction from diverse and complex spreadsheet formats.

736 REFERENCES

- 737 [1] Rui Abreu, Jácome Cunha, Joao Paulo Fernandes, Pedro Martins, Alexandre Perez, and Joao Saraiva. 2014. Faultysheet
 738 detective: When smells meet fault localization. In *2014 IEEE International Conference on Software Maintenance and*
 739 *Evolution*. IEEE, 625–628.
- 740 [2] Anonymous. [n.d.] Spreadsheet Data Extractor (SDE). https://anonymous.4open.science/r/spreadsheet_data_extractor-13BD/README.md. Accessed: 2024-12-08.
- 741 [3] Alexander Aue, Andrea Ackermann, and Norbert Röder. 2024. Converting data organised for visual perception into
 742 machine-readable formats. In *44. GIL-Jahrestagung, Biodiversität fördern durch digitale Landwirtschaft*. Gesellschaft für
 743 Informatik eV, 179–184.
- 744 [4] Daniel W Barowy, Sumit Gulwani, Ted Hart, and Benjamin Zorn. 2015. FlashRelate: extracting relational data from
 745 semi-structured spreadsheets using examples. *ACM SIGPLAN Notices* 50, 6 (2015), 218–228.
- 746 [5] D.J. Berndt, J.W. Fisher, A.R. Hevner, and J. Studnicki. 2001. Healthcare data warehousing and quality assurance.
Computer 34, 12 (2001), 56–65. <https://doi.org/10.1109/2.970578>
- 747 [6] Zhe Chen, Michael Cafarella, Jun Chen, Daniel Prevo, and Junfeng Zhuang. 2013. Senbazuru: A prototype spreadsheet
 748 database management system. *Proceedings of the VLDB Endowment* 6, 12, 1202–1205.
- 749 [7] Jácome Cunha, Joao Paulo Fernandes, Pedro Martins, Jorge Mendes, and Joao Saraiva. 2012. Smellsheet detective:
 750 A tool for detecting bad smells in spreadsheets. In *2012 IEEE Symposium on Visual Languages and Human-Centric*
Computing (VL/HCC). IEEE, 243–244.
- 751 [8] Kaval Desai. 2024. Excel Dart Package. <https://pub.dev/packages/excel>. Accessed: 2024-12-03.
- 752 [9] Haoyu Dong, Shijie Liu, Shi Han, Zhouyu Fu, and Dongmei Zhang. 2019. Tablesense: Spreadsheet table detection with
 753 convolutional neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 33, 69–76.
- 754 [10] Angus Dunn. 2010. Spreadsheets-the Good, the Bad and the Downright Ugly. *arXiv preprint arXiv:1009.5705* (2010).
- 755 [11] Julian Eberius, Christopher Werner, Maik Thiele, Katrin Braunschweig, Lars Dannecker, and Wolfgang Lehner. 2013. DeExcelerator: a framework for extracting relational data from partially structured documents. In *Proceedings of the*
22nd ACM international conference on Information & Knowledge Management. 2477–2480.
- 756 [12] Elvis Koci, Dana Kuban, Nico Luettig, Dominik Olwig, Maik Thiele, Julius Gonsior, Wolfgang Lehner, and Oscar
 757 Romero. 2019. Xlindy: Interactive recognition and information extraction in spreadsheets. In *Proceedings of the ACM*
Symposium on Document Engineering 2019. 1–4.
- 758 [13] Kate Lovett. 2023. `two_dimensional_scrollables` package - Commit 4c16f3e. <https://github.com/flutter/packages/commit/4c16f3ef40333aa0aebe8a1e46ef7b9fef9a1c1f> Accessed: 2023-08-17.
- 759 [14] Kelly Mack, John Lee, Kevin Chang, Karrie Karahalios, and Aditya Parameswaran. 2018. Characterizing scalability
 760 issues in spreadsheet software using online forums. In *Extended Abstracts of the 2018 CHI Conference on Human Factors*
in Computing Systems. 1–9.
- 761 [15] Sajjadur Rahman, Mangesh Bendre, Yuyang Liu, Shichu Zhu, Zhaoyuan Su, Karrie Karahalios, and Aditya G
 762 Parameswaran. 2021. NOAH: interactive spreadsheet exploration with dynamic hierarchical overviews. *Proceedings of*
the VLDB Endowment 14, 6 (2021), 970–983.
- 763 [16] Gursharan Singh, Leah Findlater, Kentaro Toyama, Scott Helmer, Rikin Gandhi, and Ravin Balakrishnan. 2009. Nu-
 764 meric paper forms for NGOs. In *2009 International Conference on Information and Communication Technologies and*
Development (ICTD). IEEE, 406–416.
- 765 [17] Statistisches Bundesamt (Destatis). 2024. *Statistischer Bericht - Pflegekräftevorausberechnung - 2024 bis 2070*. Technical
 766 Report. Statistisches Bundesamt, Wiesbaden, Germany. <https://www.destatis.de/DE/Themen/Gesellschaft-Umwelt/Bevoelkerung/Bevoelkerungsvorausberechnung/Publikationen/Downloads-Vorausberechnung/statistischer-bericht-pflegekraeftevorausberechnung-2070-5124210249005.html> Statistical Report - Projection of Nursing Staff - 2024 to 2070.
- 767 [18] Statistisches Bundesamt (Destatis). 2024. Statistischer Bericht: Rechnungsergebnis der Kernhaushalte der Gemeinden.
 768 <https://www.destatis.de/DE/Themen/Staat/Oeffentliche-Finanzen/Ausgaben-Einnahmen/Publikationen/Downloads-Ausgaben-und-Einnahmen/statistischer-bericht-rechnungsergebnis-kernhaushalt-gemeinden-2140331217005.html> Accessed: 2024-11-29.
- 769 [19] Alaaeddin Swidan, Felienne Hermans, and Ruben Koesoemowidjojo. 2016. Improving the performance of a large scale
 770 spreadsheet: a case study. In *2016 IEEE 23rd International Conference on Software Analysis, Evolution, and Reengineering*
 771 (*SANER*), Vol. 1. IEEE, 673–677.
- 772 [20] Dixin Tang, Fanchao Chen, Christopher De Leon, Tana Wattanawaroong, Jeaseok Yun, Srinivasan Seshadri, and Aditya G
 773 Parameswaran. 2023. Efficient and Compact Spreadsheet Formula Graphs. In *2023 IEEE 39th International Conference*
on Data Engineering (ICDE). IEEE, 2634–2646.
- 774 [21] Hannah West and Gina Green. 2008. Because excel will mind me! the state of constituent data management in small
 775 nonprofit organizations. In *Proceedings of the Fourteenth Americas Conference on Information Systems*. Association for
 776 Information Systems, AIS Electronic Library (AISel). <https://aisel.aisnet.org/amcis2008/336>

785 Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009

786

787

788

789

790

791

792

793

794

795

796

797

798

799

800

801

802

803

804

805

806

807

808

809

810

811

812

813

814

815

816

817

818

819

820

821

822

823

824

825

826

827

828

829

830

831

832

833